

ABHINAY KOTLA

+1 469 674 1021 [◇ abhinaykotla@gmail.com](mailto:abhinaykotla@gmail.com)

[LinkedIn](#) [GitHub](#) [Portfolio](#)

SUMMARY

6+ years shipping production AI systems, the most recent one from zero as founding engineer. I build agentic workflows end to end: agent logic, retrieval and context assembly, and the interface people actually use. Python and TypeScript, frontier models across six providers, and systems engineered to be correct under load, not just working in a demo. Give me a vague problem and no spec and I'll ship it. Masters in CS. Agentic coding is how I work, not a line on a skills list.

SKILLS

Languages: Python, TypeScript, JavaScript, SQL, C++, Java, C#, Bash

Agentic AI & LLMs: Agentic Workflows, Multi-Step Tool Use, MCP, LangGraph, Anthropic SDK, Frontier Model APIs (Anthropic, Gemini/Vertex, OpenAI), RAG, Embeddings & Vector Search, Context Engineering, Prompt Engineering, Model Evaluation & Selection, Fine-Tuning, Multimodal AI

AI Dev Tools: Claude Code, Cursor, GitHub Copilot, agent-assisted refactors, AI-native dev workflows

Full Stack: React, Next.js, Node.js, Express, TypeScript, REST, GraphQL, PostgreSQL, Supabase, pgvector, MongoDB, OAuth/JWT, Real-time Systems

Cloud & Reliability: AWS (S3, Lambda, EC2), Cloudflare (R2, Workers, Edge Functions), Azure, Docker, Kubernetes, CI/CD, GitHub Actions, OpenTelemetry, Monitoring & Tracing

ML Foundations: PyTorch, TensorFlow, Transformers, NLP, Knowledge Distillation, Model Compression, System Design, Distributed Systems, Performance Optimization

EXPERIENCE

Founding AI Engineer, Anthum.ai

Jan 2026 – Present

- Own the creator-facing AI Studio end to end – React/TypeScript node-graph canvas that wires 100+ models into multi-step generation workflows; shipped 160+ PRs in the first four months using agentic coding tools across the entire stack.
- Building a retrieval-grounded AI Copilot on Vertex Gemini + pgvector (768-dim embeddings): multimodal ingestion pipeline that indexes and assembles context for agent-driven ad creation, with an evaluation loop feeding ongoing fine-tuning.
- Architected provider-agnostic inference across 6 providers and 38 edge functions (FAL, Vertex, Kling, GMI Cloud) with DB-driven routing and response normalization – swap a model or a provider with zero redeploy.
- Led a five-phase pipeline re-architecture on a declarative field-mapping pattern across 10+ model adapters, cutting latency from 135s to 7s. Built the Cloudflare R2 storage layer from scratch (AWS SigV4 over Web Crypto) with parallelized transcode and archival to halve tail latency.
- Built subscription, usage-metering, and smart-throttling on row-locked RPCs that hold correct under concurrent generations – no double-spend, no abuse path – with OpenTelemetry-instrumented billing for end-to-end traceability.

Information Technology Operations Specialist, UTA Office of Information Technology

Aug 2024 – Dec 2025

- Built an AI support assistant (Python, NLP, LLM-generated responses) wired into ServiceNow; intent detection plus retrieval-grounded answers cut ticket resolution time 18% on live operational casework.

AI Engineer, Saintech Solutions LLC

Aug 2022 – Dec 2023

- First technical hire; took AI features and full-stack systems from zero to production – web apps, backend services, APIs, and the data pipelines behind them.

Full Stack Engineer, HA Talent Tech

Jan 2021 – Aug 2022

- Built and deployed a production full-stack app (React, Node.js, REST) with secure auth and RBAC, sustaining 99.9% uptime.

Machine Learning Engineer, SmartKnower

Jul 2020 – Dec 2020

- Built ML pipelines and trained classification models for adaptive course recommendations serving 10,000+ learners; lifted engagement 23% and cut inference latency 40%.

ML Engineer, 1StopAI

Jan 2020 – Jun 2020

- Built and deployed emotion and gender detection models on support audio, tuning inference pipelines for real-time routing.

PROJECTS

News Summarization with T5

LLMs, NLP, Evaluation

- Fine-tuned a T5 transformer for abstractive summarization to ROUGE-1 0.53, building the full pipeline: preprocessing, decoding-strategy control, and automated evaluation to decide which outputs were good enough to ship.

FinAI Personal Finance Assistant

LLM Systems, Live Data

- Finance assistant grounded in live Plaid account data; assembled third-party financial context and real-time market sentiment into LLM-generated guidance.

Efficient CV Models with Knowledge Distillation

Model Optimization

- Compressed a CNN by 99% while holding 97% accuracy via structured pruning and teacher-student distillation, with FP16 quantization for a 10× GPU-memory reduction.

EDUCATION

Masters in Computer Science

University of Texas at Arlington | GPA: 4.0/4.0

BE in Computer Science

GITAM, Hyderabad | CGPA: 8.47

Certifications

Neural Networks and Deep Learning (DeepLearning.AI)

Robotic Process Automation (RPA) – IBM

Big Data (Spark, Hadoop)